

実践演習9-2

IMDBデータは映画のレビューに対して、P/N(肯定/否定)のラベルが付いた学習データです。学習用に25000事例、評価用に25000事例用意されていて、PNの割合はそれぞれ50%です。各レビューは単語列ではなく、単語インデックスの系列として表現されています。

ここでは、頻度上位10000語を対象とし、データの大きさは先頭の50単語に限定します。

In [0]:

```
import tensorflow as tf
from tensorflow import keras
```

In [2]:

```
max_features = 10000
maxlen = 50
(X_train, y_train), (X_test, y_test) = keras.datasets.imdb.load_data(num_words=max_features)
X_train = keras.preprocessing.sequence.pad_sequences(X_train, maxlen=maxlen)
X_test = keras.preprocessing.sequence.pad_sequences(X_test, maxlen=maxlen)
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/imdb.npz> (<https://storage.googleapis.com/tensorflow/tf-keras-datasets/imdb.npz>)
17465344/17464789 [=====] - 1s 0us/step

単語インデックスを単語に戻して、元のデータを確認します。インデックスは"padding", "start of sequence", "unknown"にそれぞれ0,1,2が割り当てられているので、3つずらして対応させます。

In [3]:

```
word_index = keras.datasets.imdb.get_word_index()
reverse_word_index = dict([(value, key) for (key, value) in word_index.items()])
decoded_review = ' '.join([reverse_word_index.get(i - 3, '?') for i in X_train[0]])
decoded_review
```

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/imdb_word_index.json (https://storage.googleapis.com/tensorflow/tf-keras-datasets/imdb_word_index.json)
1646592/1641221 [=====] - 0s 0us/step

Out[3]:

"grown up are such a big profile for the whole film but these children are amazing and should be praised for what they have done don't you think the whole story was so lovely because it was true and was someone's life after all that was shared with us all"

単純なRNNを構成して学習させます。

In [4]:

```
model = keras.Sequential([
    keras.layers.Embedding(max_features, 128),
    keras.layers.SimpleRNN(64),
    keras.layers.Dense(1, activation='sigmoid')
])
model.summary()
```

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, None, 128)	1280000
simple_rnn (SimpleRNN)	(None, 64)	12352
dense (Dense)	(None, 1)	65
Total params: 1,292,417		
Trainable params: 1,292,417		
Non-trainable params: 0		

In [5]:

```
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['acc'])
model.fit(X_train, y_train, epochs=5, batch_size=200, validation_split=0.2)
```

/usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/gradients_impl.py:112:
 UserWarning: Converting sparse IndexedSlices to a dense Tensor of unknown shape. This
 may consume a large amount of memory.
 "Converting sparse IndexedSlices to a dense Tensor of unknown shape. "

Train on 20000 samples, validate on 5000 samples

Epoch 1/5

20000/20000 [=====] - 5s 232us/step - loss: 0.5202 - acc:
 0.7340 - val_loss: 0.4436 - val_acc: 0.7994

Epoch 2/5

20000/20000 [=====] - 4s 181us/step - loss: 0.2958 - acc:
 0.8796 - val_loss: 0.4864 - val_acc: 0.7850

Epoch 3/5

20000/20000 [=====] - 4s 181us/step - loss: 0.1255 - acc:
 0.9585 - val_loss: 0.5739 - val_acc: 0.7730

Epoch 4/5

20000/20000 [=====] - 4s 185us/step - loss: 0.0308 - acc:
 0.9934 - val_loss: 0.7463 - val_acc: 0.7634

Epoch 5/5

20000/20000 [=====] - 4s 180us/step - loss: 0.0065 - acc:
 0.9995 - val_loss: 0.8327 - val_acc: 0.7730

Out[5]:

<tensorflow.python.keras.callbacks.History at 0x7f5378f495c0>

In [6]:

```
score = model.evaluate(X_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

Test loss: 0.7796998586845398

Test accuracy: 0.77572

RNNユニットをLSTMに変更して性能の変化を確認します。

In [7]:

```

model = keras.Sequential([
    keras.layers.Embedding(max_features, 128),
    keras.layers.LSTM(64),
    keras.layers.Dense(1, activation=tf.nn.sigmoid)
])
model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['acc'])
model.fit(X_train, y_train, epochs=5, batch_size=200, validation_split=0.2)

```

/usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/gradients_impl.py:112:
 UserWarning: Converting sparse IndexedSlices to a dense Tensor of unknown shape. This
 may consume a large amount of memory.

"Converting sparse IndexedSlices to a dense Tensor of unknown shape. "

Train on 20000 samples, validate on 5000 samples

Epoch 1/5

20000/20000 [=====] - 15s 745us/step - loss: 0.5037 - acc:
 0.7482 - val_loss: 0.4238 - val_acc: 0.8050

Epoch 2/5

20000/20000 [=====] - 14s 714us/step - loss: 0.3452 - acc:
 0.8500 - val_loss: 0.4239 - val_acc: 0.8024

Epoch 3/5

20000/20000 [=====] - 15s 728us/step - loss: 0.2997 - acc:
 0.8757 - val_loss: 0.4600 - val_acc: 0.8034

Epoch 4/5

20000/20000 [=====] - 15s 734us/step - loss: 0.2653 - acc:
 0.8935 - val_loss: 0.4701 - val_acc: 0.8030

Epoch 5/5

20000/20000 [=====] - 15s 730us/step - loss: 0.2353 - acc:
 0.9081 - val_loss: 0.5012 - val_acc: 0.7924

Out[7]:

<tensorflow.python.keras.callbacks.History at 0x7f53787a45c0>

In [8]:

```

score = model.evaluate(X_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

Test loss: 0.49411070870399476

Test accuracy: 0.79984

双方向型にして、ユニットをGRUにします。また、ドロップアウトも加えてみます。

In [9]:

```

model = keras.Sequential([
    keras.layers.Embedding(max_features, 64),
    keras.layers.Bidirectional(keras.layers.GRU(32)),
    keras.layers.Dropout(0.5),
    keras.layers.Dense(1, activation=tf.nn.sigmoid)
])
model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['acc'])
model.fit(X_train, y_train, epochs=5, batch_size=200, validation_split=0.2)

```

/usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/gradients_impl.py:112:
 UserWarning: Converting sparse IndexedSlices to a dense Tensor of unknown shape. This
 may consume a large amount of memory.

"Converting sparse IndexedSlices to a dense Tensor of unknown shape. "

Train on 20000 samples, validate on 5000 samples

Epoch 1/5

20000/20000 [=====] - 24s 1ms/step - loss: 0.5917 - acc: 0.6724 - val_loss: 0.4531 - val_acc: 0.7824

Epoch 2/5

20000/20000 [=====] - 23s 1ms/step - loss: 0.3716 - acc: 0.8401 - val_loss: 0.4077 - val_acc: 0.8116

Epoch 3/5

20000/20000 [=====] - 22s 1ms/step - loss: 0.3097 - acc: 0.8701 - val_loss: 0.4608 - val_acc: 0.8050

Epoch 4/5

20000/20000 [=====] - 21s 1ms/step - loss: 0.2695 - acc: 0.8899 - val_loss: 0.4381 - val_acc: 0.8048

Epoch 5/5

20000/20000 [=====] - 21s 1ms/step - loss: 0.2395 - acc: 0.9055 - val_loss: 0.4686 - val_acc: 0.7960

Out[9]:

<tensorflow.python.keras.callbacks.History at 0x7f5371428f28>

In [10]:

```

score = model.evaluate(X_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

Test loss: 0.4524161660003662

Test accuracy: 0.80304

In [0]:

